

1. CUSTOMER DATA

 3.2.1 Jobs Stages Storage Environment Executors SQL ex1_customer_data application UI

Spark Jobs (?)

User: racit
Total Uptime: 1.5 min
Scheduling Mode: FIFO
Completed Jobs: 6

► Event Timeline
▼ Completed Jobs (6)

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5	parquet at ex1_customer_data.scala:40 parquet at ex1_customer_data.scala:40	2025/12/02 23:15:18	1 s	1/1 (1 skipped)	1/1 (50 skipped)
4	parquet at ex1_customer_data.scala:40 parquet at ex1_customer_data.scala:40	2025/12/02 23:15:09	10 s	1/1	50/50
3	json at ex1_customer_data.scala:39 json at ex1_customer_data.scala:39	2025/12/02 23:15:08	0.2 s	1/1 (1 skipped)	1/1 (50 skipped)
2	json at ex1_customer_data.scala:39 json at ex1_customer_data.scala:39	2025/12/02 23:14:58	10 s	1/1	50/50
1	csv at ex1_customer_data.scala:38 csv at ex1_customer_data.scala:38	2025/12/02 23:14:58	0.2 s	1/1 (1 skipped)	1/1 (50 skipped)
0	csv at ex1_customer_data.scala:38 csv at ex1_customer_data.scala:38	2025/12/02 23:14:46	11 s	1/1	50/50

Stages for All Jobs

Completed Stages: 6
Skipped Stages: 3

▼ Completed Stages (6)

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
8	parquet at ex1_customer_data.scala:40	+details 2025/12/02 23:15:18	1 s	1/1	1121.0 B	159.7 KIB		
6	parquet at ex1_customer_data.scala:40	+details 2025/12/02 23:15:09	10 s	50/50			159.7 KIB	
5	json at ex1_customer_data.scala:39	+details 2025/12/02 23:15:08	0.2 s	1/1	1664.0 B	159.7 KIB		
3	json at ex1_customer_data.scala:39	+details 2025/12/02 23:14:58	10 s	50/50			159.7 KIB	
2	csv at ex1_customer_data.scala:38	+details 2025/12/02 23:14:58	0.2 s	1/1	713.0 B	159.7 KIB		
0	csv at ex1_customer_data.scala:38	+details 2025/12/02 23:14:46	11 s	50/50			159.7 KIB	

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

▼ Skipped Stages (3)

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
7	parquet at ex1_customer_data.scala:40	+details Unknown	Unknown	0/50				
4	json at ex1_customer_data.scala:39	+details Unknown	Unknown	0/50				
1	csv at ex1_customer_data.scala:38	+details Unknown	Unknown	0/50				

Executors

► Show Additional Metrics

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	83.3 KIB / 434.4 MiB	0.0 B	8	0	0	153	153	4.1 min (5 s)	0.0 B	479.1 KIB	479.1 KIB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	83.3 KIB / 434.4 MiB	0.0 B	8	0	0	153	153	4.1 min (5 s)	0.0 B	479.1 KIB	479.1 KIB	0

Executors

Show 20 entries Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.1.12:61492	Active	0	83.3 KIB / 434.4 MiB	0.0 B	8	0	0	153	153	4.1 min (5 s)	0.0 B	479.1 KIB	479.1 KIB	Thread Dump

Showing 1 to 1 of 1 entries Previous 1 Next

2. SALES DATA

Spark 3.2.1 Jobs Stages Storage Environment Executors SQL ex2_sales_data application UI

Spark Jobs (2)

User: racit
Total Uptime: 38 s
Scheduling Mode: FIFO
Completed Jobs: 6

► Event Timeline

▼ Completed Jobs (6)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5	json at ex2_sales_data.scala:40 json at ex2_sales_data.scala:40	2025/12/02 23:38:37	0.1 s	1/1 (1 skipped)	1/1 (50 skipped)
4	json at ex2_sales_data.scala:40 json at ex2_sales_data.scala:40	2025/12/02 23:38:34	3 s	1/1	50/50
3	csv at ex2_sales_data.scala:39 csv at ex2_sales_data.scala:39	2025/12/02 23:38:34	0.1 s	1/1 (1 skipped)	1/1 (50 skipped)
2	csv at ex2_sales_data.scala:39 csv at ex2_sales_data.scala:39	2025/12/02 23:38:31	2 s	1/1	50/50
1	parquet at ex2_sales_data.scala:38 parquet at ex2_sales_data.scala:38	2025/12/02 23:38:31	0.7 s	1/1 (1 skipped)	1/1 (50 skipped)
0	parquet at ex2_sales_data.scala:38 parquet at ex2_sales_data.scala:38	2025/12/02 23:38:27	4 s	1/1	50/50

Stages for All Jobs

Completed Stages: 6

Skipped Stages: 3

▼ Completed Stages (6)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
8	json at ex2_sales_data.scala:40 +details	2025/12/02 23:38:37	89 ms	1/1	5.6 KiB	310.0 KiB		
6	json at ex2_sales_data.scala:40 +details	2025/12/02 23:38:34	3 s	50/50			310.0 KiB	
5	csv at ex2_sales_data.scala:39 +details	2025/12/02 23:38:34	98 ms	1/1	2.9 KiB	310.0 KiB		
3	csv at ex2_sales_data.scala:39 +details	2025/12/02 23:38:31	2 s	50/50			310.0 KiB	
2	parquet at ex2_sales_data.scala:38 +details	2025/12/02 23:38:31	0.7 s	1/1	1971.0 B	310.0 KiB		
0	parquet at ex2_sales_data.scala:38 +details	2025/12/02 23:38:27	4 s	50/50			310.0 KiB	

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

▼ Skipped Stages (3)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
7	json at ex2_sales_data.scala:40 +details	Unknown	Unknown	0/50				
4	csv at ex2_sales_data.scala:39 +details	Unknown	Unknown	0/50				
1	parquet at ex2_sales_data.scala:38 +details	Unknown	Unknown	0/50				

Executors

► Show Additional Metrics

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	85.3 KiB / 434.4 MiB	0.0 B	8	0	0	153	153	1.2 min (1 s)	0.0 B	930 KiB	930 KiB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	85.3 KiB / 434.4 MiB	0.0 B	8	0	0	153	153	1.2 min (1 s)	0.0 B	930 KiB	930 KiB	0

Executors

Show 20 entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.1.12:61670	Active	0	85.3 KiB / 434.4 MiB	0.0 B	8	0	0	153	153	1.2 min (1 s)	0.0 B	930 KiB	930 KiB	Thread Dump

Showing 1 to 1 of 1 entries

Previous 1 Next

3. LOG FILES

Spark Jobs [\(?\)](#)

User: racit
 Total Uptime: 44 s
 Scheduling Mode: FIFO
 Active Jobs: 1
 Completed Jobs: 3

[Event Timeline](#)

Active Jobs (1)

Page: 1

1 Pages. Jump to . Show items in a page. [Go](#)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
3	runJob at SparkHadoopWriter.scala:83 runJob at SparkHadoopWriter.scala:83 (kill)	2025/12/02 23:49:27	4 s	0/1	7/40 (9 running)

Page: 1

1 Pages. Jump to . Show items in a page. [Go](#)

Completed Jobs (3)

Page: 1

1 Pages. Jump to . Show items in a page. [Go](#)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
2	count at ex3_log_files.scala:37 count at ex3_log_files.scala:37	2025/12/02 23:49:27	63 ms	1/1 (1 skipped)	1/1 (40 skipped)
1	count at ex3_log_files.scala:37 count at ex3_log_files.scala:37	2025/12/02 23:49:10	17 s	1/1	40/40
0	count at ex3_log_files.scala:34 count at ex3_log_files.scala:34	2025/12/02 23:48:50	20 s	1/1	40/40

 3.2.1 Jobs Stages Storage Environment Executors SQL ex3_log_files application UI

Stages for All Jobs

Completed Stages: 5
 Skipped Stages: 1

Completed Stages (5)

Page: 1

1 Pages. Jump to . Show items in a page. [Go](#)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
5	json at ex3_log_files.scala:44 +details	2025/12/02 23:49:43	16 s	40/40	425.6 MiB			
4	runJob at SparkHadoopWriter.scala:83 +details	2025/12/02 23:49:27	15 s	40/40	65.5 MiB			
3	count at ex3_log_files.scala:37 +details	2025/12/02 23:49:27	45 ms	1/1		2.3 KiB		
1	count at ex3_log_files.scala:37 +details	2025/12/02 23:49:10	17 s	40/40			2.3 KiB	
0	count at ex3_log_files.scala:34 +details	2025/12/02 23:48:50	20 s	40/40				

Page: 1

1 Pages. Jump to . Show items in a page. [Go](#)

Skipped Stages (1)

Page: 1

1 Pages. Jump to . Show items in a page. [Go](#)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	count at ex3_log_files.scala:37 +details	Unknown	Unknown	0/40				

Page: 1

1 Pages. Jump to . Show items in a page. [Go](#)

 3.2.1 Jobs Stages Storage Environment Executors SQL ex3_log_files application UI

Executors

[Show Additional Metrics](#)

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	120.7 KiB / 434.4 MiB	0.0 B	8	0	0	161	161	9.0 min (10 s)	0.0 B	2.3 KiB	2.3 KiB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	120.7 KiB / 434.4 MiB	0.0 B	8	0	0	161	161	9.0 min (10 s)	0.0 B	2.3 KiB	2.3 KiB	0

Executors

Show 20 entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.1.12:61721	Active	0	120.7 KiB / 434.4 MiB	0.0 B	8	0	0	161	161	9.0 min (10 s)	0.0 B	2.3 KiB	2.3 KiB	Thread Dump

Showing 1 to 1 of 1 entries

Previous [1](#) Next

PRODUCT CATALOG

Spark Jobs [\(2\)](#)

User: racit

Total Uptime: 1.7 min

Scheduling Mode: FIFO

Completed Jobs: 6

[▶ Event Timeline](#)

[▼ Completed Jobs \(6\)](#)

Page:

1 Pages. Jump to . Show items in a page. [Go](#)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5	parquet at ex4_product_catalog.scala:42 parquet at ex4_product_catalog.scala:42	2025/12/02 23:53:16	2 s	1/1 (1 skipped)	1/1 (50 skipped)
4	parquet at ex4_product_catalog.scala:42 parquet at ex4_product_catalog.scala:42	2025/12/02 23:53:13	3 s	1/1	50/50
3	json at ex4_product_catalog.scala:41 json at ex4_product_catalog.scala:41	2025/12/02 23:53:12	0.1 s	1/1 (1 skipped)	1/1 (50 skipped)
2	json at ex4_product_catalog.scala:41 json at ex4_product_catalog.scala:41	2025/12/02 23:53:09	3 s	1/1	50/50
1	csv at ex4_product_catalog.scala:40 csv at ex4_product_catalog.scala:40	2025/12/02 23:53:09	0.2 s	1/1 (1 skipped)	1/1 (50 skipped)
0	csv at ex4_product_catalog.scala:40 csv at ex4_product_catalog.scala:40	2025/12/02 23:53:05	4 s	1/1	50/50

Stages for All Jobs

Completed Stages: 6

Skipped Stages: 3

[▼ Completed Stages \(6\)](#)

Page:

1 Pages. Jump to . Show items in a page. [Go](#)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
8	parquet at ex4_product_catalog.scala:42 +details	2025/12/02 23:53:16	2 s	1/1	1116.0 B	24.8 KIB		
6	parquet at ex4_product_catalog.scala:42 +details	2025/12/02 23:53:13	3 s	50/50				24.8 KIB
5	json at ex4_product_catalog.scala:41 +details	2025/12/02 23:53:12	80 ms	1/1	402.0 B	24.8 KIB		
3	json at ex4_product_catalog.scala:41 +details	2025/12/02 23:53:09	3 s	50/50				24.8 KIB
2	csv at ex4_product_catalog.scala:40 +details	2025/12/02 23:53:09	0.1 s	1/1	198.0 B	24.8 KIB		
0	csv at ex4_product_catalog.scala:40 +details	2025/12/02 23:53:05	4 s	50/50				24.8 KIB

Page:

1 Pages. Jump to . Show items in a page. [Go](#)

[▼ Skipped Stages \(3\)](#)

Page:

1 Pages. Jump to . Show items in a page. [Go](#)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
7	parquet at ex4_product_catalog.scala:42 +details	Unknown	Unknown	0/50				
4	json at ex4_product_catalog.scala:41 +details	Unknown	Unknown	0/50				
1	csv at ex4_product_catalog.scala:40 +details	Unknown	Unknown	0/50				

 3.2.1 Jobs Stages Storage Environment Executors SQL ex4_product_catalog application UI

Executors

[▶ Show Additional Metrics](#)

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	104.7 KIB / 434.4 MiB	0.0 B	8	0	0	153	153	1.4 min (2 s)	0.0 B	74.3 KIB	74.3 KIB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	104.7 KIB / 434.4 MiB	0.0 B	8	0	0	153	153	1.4 min (2 s)	0.0 B	74.3 KIB	74.3 KIB	0

Executors

Show entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.12.61:757	Active	0	104.7 KIB / 434.4 MiB	0.0 B	8	0	0	153	153	1.4 min (2 s)	0.0 B	74.3 KIB	74.3 KIB	Thread Dump

Showing 1 to 1 of 1 entries

Previous Next

5. IoT SENSOR DATA

 3.2.1	Jobs	Stages	Storage	Environment	Executors	SQL	ex5_iot_sensor application UI
---	------	--------	---------	-------------	-----------	-----	-------------------------------

Spark Jobs (?)

User: racit
 Total Uptime: 30 s
 Scheduling Mode: FIFO
 Completed Jobs: 4

► Event Timeline

▼ Completed Jobs (4)

Page: 1

1 Pages. Jump to . Show items in a page.

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
3	json at ex5_iot_sensor.scala:41 json at ex5_iot_sensor.scala:41	2025/12/02 23:57:01	0.2 s	1/1 (1 skipped)	1/1 (0 skipped)
2	json at ex5_iot_sensor.scala:41 json at ex5_iot_sensor.scala:41	2025/12/02 23:56:56	5 s	1/1	50/50
1	parquet at ex5_iot_sensor.scala:40 parquet at ex5_iot_sensor.scala:40	2025/12/02 23:56:55	0.9 s	1/1 (1 skipped)	1/1 (0 skipped)
0	parquet at ex5_iot_sensor.scala:40 parquet at ex5_iot_sensor.scala:40	2025/12/02 23:56:47	8 s	1/1	50/50

Page: 1

1 Pages. Jump to . Show items in a page.

 3.2.1	Jobs	Stages	Storage	Environment	Executors	SQL	ex5_iot_sensor application UI
---	------	--------	---------	-------------	-----------	-----	-------------------------------

Stages for All Jobs

Completed Stages: 4
 Skipped Stages: 2

▼ Completed Stages (4)

Page: 1

1 Pages. Jump to . Show items in a page.

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
5	json at ex5_iot_sensor.scala:41 +details	2025/12/02 23:57:01	0.1 s	1/1	8.4 Kib	431.8 Kib		
3	json at ex5_iot_sensor.scala:41 +details	2025/12/02 23:56:56	5 s	50/50			431.8 Kib	
2	parquet at ex5_iot_sensor.scala:40 +details	2025/12/02 23:56:55	0.8 s	1/1	3.0 Kib	432.0 Kib		
0	parquet at ex5_iot_sensor.scala:40 +details	2025/12/02 23:56:47	8 s	50/50			432.0 Kib	

Page: 1

1 Pages. Jump to . Show items in a page.

▼ Skipped Stages (2)

Page: 1

1 Pages. Jump to . Show items in a page.

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
4	json at ex5_iot_sensor.scala:41 +details	Unknown	Unknown	0/50				
1	parquet at ex5_iot_sensor.scala:40 +details	Unknown	Unknown	0/50				

Page: 1

1 Pages. Jump to . Show items in a page.

 3.2.1	Jobs	Stages	Storage	Environment	Executors	SQL	ex5_iot_sensor application UI
---	------	--------	---------	-------------	-----------	-----	-------------------------------

Executors

► Show Additional Metrics

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	39.3 Kib / 434.4 MiB	0.0 B	8	0	0	102	102	1.7 min (4 s)	0.0 B	863.7 Kib	863.7 Kib	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	39.3 Kib / 434.4 MiB	0.0 B	8	0	0	102	102	1.7 min (4 s)	0.0 B	863.7 Kib	863.7 Kib	0

Executors

Show entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.1.12:61787	Active	0	39.3 Kib / 434.4 MiB	0.0 B	8	0	0	102	1.7 min (4 s)	0.0 B	863.7 Kib	863.7 Kib	Thread Dump	

Showing 1 to 1 of 1 entries

Previous Next

6. SOCIAL MEDIA USERS + POSTS

Spark 3.2.1 Jobs Stages Storage Environment Executors SQL ex6_social_media application UI

Spark Jobs (?)

User: racit
Total Uptime: 12 s
Scheduling Mode: FIFO
Active Jobs: 1
Completed Jobs: 1

[Event Timeline](#)

Active Jobs (1)

Page: 1

1 Pages. Jump to . Show items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	json at ex6_social_media.scala:57 json at ex6_social_media.scala:57 (kill)	2025/12/03 00:38:53	9 s	0/1	20/30 (9 running)

Page: 1

1 Pages. Jump to . Show items in a page. Go

Completed Jobs (1)

Page: 1

1 Pages. Jump to . Show items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	json at ex6_social_media.scala:57 json at ex6_social_media.scala:57	2025/12/03 00:38:53	7 s	1/1	40/40

Jobs Stages Storage Environment Executors SQL ex6_social_media application UI

Stages for All Jobs

Completed Stages: 6
Skipped Stages: 12

Completed Stages (6)

Page: 1

1 Pages. Jump to . Show items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
17	json at ex6_social_media.scala:57	+details 2025/12/03 00:39:05	0.1 s	1/1	1740.0 B	3.7 kB		
12	json at ex6_social_media.scala:57	+details 2025/12/03 00:39:05	0.1 s	1/1		26.1 kB	3.7 kB	
8	json at ex6_social_media.scala:57	+details 2025/12/03 00:39:05	35 ms	1/1		26.1 kB		
4	json at ex6_social_media.scala:57	+details 2025/12/03 00:39:03	2 s	8/8		20.7 MiB	26.1 kB	
1	json at ex6_social_media.scala:57	+details 2025/12/03 00:38:53	3 s	30/30			10.4 MiB	
0	json at ex6_social_media.scala:57	+details 2025/12/03 00:38:53	7 s	40/40			10.3 MiB	

Page: 1

1 Pages. Jump to . Show items in a page. Go

Skipped Stages (12)

Page: 1

1 Pages. Jump to . Show items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
16	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/1				
15	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/8				
14	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/40				
13	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/30				
11	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/8				
10	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/40				
9	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/30				
7	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/8				
6	json at ex6_social_media.scala:57	+details Unknown	Unknown	0/40				

Jobs Stages Storage Environment Executors SQL ex6_social_media application UI

Executors

[Show Additional Metrics](#)

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	0.0 B / 434.4 MiB	0.0 B	8	0	0	81	81	1.4 min (6 s)	0.0 B	20.8 MiB	20.8 MiB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	0.0 B / 434.4 MiB	0.0 B	8	0	0	81	81	1.4 min (6 s)	0.0 B	20.8 MiB	20.8 MiB	0

Executors

Show 20 entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	10.99.210.226:62038	Active	0	0.0 B / 434.4 MiB	0.0 B	8	0	0	81	81	1.4 min (6 s)	0.0 B	20.8 MiB	20.8 MiB	Thread Dump

Showing 1 to 1 of 1 entries

Previous Next

7. FINANCIAL TRANSACTIONS

Spark Jobs (1)

User: racit
Total Uptime: 1.4 min
Scheduling Mode: FIFO
Completed Jobs: 4

Event Timeline

Completed Jobs (4)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
3	json at ex7_financial_txns.scala:36 json at ex7_financial_txns.scala:36	2025/12/03 01:18:48	0.4 s	1/1 (1 skipped)	9/9 (50 skipped)
2	json at ex7_financial_txns.scala:36 json at ex7_financial_txns.scala:36	2025/12/03 01:18:42	6 s	1/1	50/50
1	parquet at ex7_financial_txns.scala:35 parquet at ex7_financial_txns.scala:35	2025/12/03 01:18:40	2 s	1/1 (1 skipped)	9/9 (50 skipped)
0	parquet at ex7_financial_txns.scala:35 parquet at ex7_financial_txns.scala:35	2025/12/03 01:18:32	8 s	1/1	50/50

Pages: 1 | 1 Pages, Jump to: 1 . Show 100 items in a page. Go

Stages for All Jobs

Completed Stages: 4
Skipped Stages: 2

Completed Stages (4)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
5	json at ex7_financial_txns.scala:36	+details 2025/12/03 01:18:48	0.3 s	9/9	4.7 MiB	39.4 MiB		
3	json at ex7_financial_txns.scala:36	+details 2025/12/03 01:18:42	6 s	50/50			39.4 MiB	
2	parquet at ex7_financial_txns.scala:35	+details 2025/12/03 01:18:40	2 s	9/9	990.6 KiB	39.5 MiB		
0	parquet at ex7_financial_txns.scala:35	+details 2025/12/03 01:18:32	8 s	50/50			39.5 MiB	

Pages: 1 | 1 Pages, Jump to: 1 . Show 100 items in a page. Go

Skipped Stages (2)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
4	json at ex7_financial_txns.scala:36	+details Unknown	Unknown	0/50				
1	parquet at ex7_financial_txns.scala:35	+details Unknown	Unknown	0/50				

Pages: 1 | 1 Pages, Jump to: 1 . Show 100 items in a page. Go

Executors

Show Additional Metrics

Summary

RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	883.7 KiB / 434.4 MiB	0.0 B	8	0	0	118	118	2.0 min (5 s)	0.0 B	78.9 MiB	78.9 MiB	0
Dead(0)	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	883.7 KiB / 434.4 MiB	0.0 B	8	0	0	118	118	2.0 min (5 s)	0.0 B	78.9 MiB	78.9 MiB	0

Executors

Show 20 entries

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.1.12:62193	Active	0	883.7 KiB / 434.4 MiB	0.0 B	8	0	0	118	118	2.0 min (5 s)	0.0 B	78.9 MiB	78.9 MiB	Thread Dump

Showing 1 to 1 of 1 entries

Search:

Previous 1 Next

8. EMPLOYEE DATA

Spark Jobs (2)

User: racit
Total Uptime: 6.1 min
Scheduling Mode: FIFO
Completed Jobs: 2

Event Timeline
Completed Jobs (2)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	csv at ex8_employee_data.scala:30 csv at ex8_employee_data.scala:30	2025/12/03 01:25:45	0.2 s	1/1 (1 skipped)	1/1 (30 skipped)
0	csv at ex8_employee_data.scala:30 csv at ex8_employee_data.scala:30	2025/12/03 01:25:43	2 s	1/1	30/30

Pages: 1 | 1 Pages. Jump to: 1 | Show 100 items in a page. Go

Stages for All Jobs

Completed Stages: 2
Skipped Stages: 1

Completed Stages (2)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	csv at ex8_employee_data.scala:30	+details 2025/12/03 01:25:45	0.2 s	1/1	86.0 B	9.0 KIB		
0	csv at ex8_employee_data.scala:30	+details 2025/12/03 01:25:43	2 s	30/30			9.0 KIB	

Pages: 1 | 1 Pages. Jump to: 1 | Show 100 items in a page. Go

- Skipped Stages (1)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	csv at ex8_employee_data.scala:30	+details Unknown	Unknown	0/30				

Pages: 1 | 1 Pages. Jump to: 1 | Show 100 items in a page. Go

Executors

Show Additional Metrics

Summary

Active(1)	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
0	103.9 KIB / 434.4 MiB	0.0 B	8	0	0	31	31	13 s (0.5 s)	0.0 B	9 KIB	9 KIB	0	
Dead(0)	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0	
Total(1)	103.9 KIB / 434.4 MiB	0.0 B	8	0	0	31	31	13 s (0.5 s)	0.0 B	9 KIB	9 KIB	0	

Executors

Show 20 entries

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.1.12:62226	Active	0	103.9 KIB / 434.4 MiB	0.0 B	8	0	0	31	31	13 s (0.5 s)	0.0 B	9 KIB	9 KIB	Thread Dump

Showing 1 to 1 of 1 entries

Search:

Previous 1 Next

9. STUDENT SCORES

Spark Jobs (2)

User: root
Total Uptime: 41 s
Scheduling Mode: FIFO
Completed Jobs: 2

- Event Timeline
- Completed Jobs (2)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	json at ex9_student_scores.scala:33 json at ex9_student_scores.scala:33	2025/12/03 01:39:00	0.5 s	1/1	20/20
0	parquet at ex9_student_scores.scala:32 parquet at ex9_student_scores.scala:32	2025/12/03 01:38:58	2 s	1/1	20/20

Page: 1 / 1 Pages, Jump to 1 Show 100 items in a page Go

Stages for All Jobs

Completed Stages: 2

- Completed Stages (2)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	json at ex9_student_scores.scala:33	+details 2025/12/03 01:39:00	0.5 s	20/20	11.7 MB			
0	parquet at ex9_student_scores.scala:32	+details 2025/12/03 01:38:58	2 s	20/20	1804.5 kB			

Page: 1 / 1 Pages, Jump to 1 Show 100 items in a page Go

Executors

Show Additional Metrics

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	74.6 KiB / 434.4 MiB	0.0 B	8	0	0	40	40	15 s (0.4 s)	0.0 B	0.0 B	0.0 B	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	74.6 KiB / 434.4 MiB	0.0 B	8	0	0	40	40	15 s (0.4 s)	0.0 B	0.0 B	0.0 B	0

Executors

Show 20 entries

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.1.12:62275	Active	0	74.6 KiB / 434.4 MiB	0.0 B	8	0	0	40	40	15 s (0.4 s)	0.0 B	0.0 B	0.0 B	Thread Dump

Showing 1 to 1 of 1 entries

Search:

Previous 1 Next

10. MINI CASE STUDY

Spark Jobs (?)

User: racit
 Total Uptime: 22 s
 Scheduling Mode: FIFO
 Active Jobs: 2
 Completed Jobs: 2

► Event Timeline

- Active Jobs (2)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
3	parquet at ex10_case_study.scala:74 parquet at ex10_case_study.scala:74	2025/12/03 01:53:46 (kill)	12 s	0/1	0/50
2	parquet at ex10_case_study.scala:74 parquet at ex10_case_study.scala:74	2025/12/03 01:53:46 (kill)	12 s	0/1	85/80 (9 running)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

- Completed Jobs (2)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	csv at ex10_case_study.scala:54 csv at ex10_case_study.scala:54	2025/12/03 01:53:43	3 s	1/1	80/80
0	csv at ex10_case_study.scala:34 csv at ex10_case_study.scala:34	2025/12/03 01:53:40	3 s	1/1	50/50

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Stages for All Jobs

Active Stages: 2
 Completed Stages: 5
 Skipped Stages: 2

- Active Stages (2)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
8	show at ex10_case_study.scala:80	+details (kill) 2025/12/03 01:54:28	Unknown	0/50				
7	show at ex10_case_study.scala:80	+details (kill) 2025/12/03 01:54:28	13 s	60/80 (9 running)				51.9 MB

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

- Completed Stages (5)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
6	parquet at ex10_case_study.scala:74	+details 2025/12/03 01:54:11	15 s	9/9	38.6 MB	98.1 MB		
3	parquet at ex10_case_study.scala:74	+details 2025/12/03 01:53:46	9 s	50/50				28.9 MB
2	parquet at ex10_case_study.scala:74	+details 2025/12/03 01:53:46	16 s	80/80				69.2 MB
1	csv at ex10_case_study.scala:54	+details 2025/12/03 01:53:43	3 s	80/80	159.2 MB			
0	csv at ex10_case_study.scala:34	+details 2025/12/03 01:53:40	3 s	50/50	31.4 MB			

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

- Skipped Stages (2)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
5	parquet at ex10_case_study.scala:74	+details Unknown	Unknown	0/80				
4	parquet at ex10_case_study.scala:74	+details Unknown	Unknown	0/50				

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Executors

► Show Additional Metrics

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	13.1 KIB / 434.4 MB	0.0 B	8	13	0	344	357	7.8 min (27 s)	0.0 B	98.1 MB	162.7 MB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	13.1 KIB / 434.4 MB	0.0 B	8	13	0	344	357	7.8 min (27 s)	0.0 B	98.1 MB	162.7 MB	0

Executors

Show 20 entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.1.12:62327	Active	0	13.1 KIB / 434.4 MB	0.0 B	8	13	0	344	357	7.8 min (27 s)	0.0 B	98.1 MB	162.7 MB	Thread Dump

Showing 1 to 1 of 1 entries

Previous 1 Next